

Bandits

Ana Matran-Fernandez
University of Essex

February 11, 2019

Introduction

Bandits

Non-stationary regimes

The adversarial case

Contextual Bandits

Conclusion

BANDITS

- ▶ We are in effect revisiting some ideas from lecture 2
 - ▶ Hypothesis testing
- ▶ This is a much easier framework to understand than hypothesis testing
- ▶ Bandits are the simplest type of reinforcement learning problem

NOT SUPERVISED LEARNING


This is **NOT** supervised learning:

- ▶ In SL we learn from samples provided by a knowledgeable supervisor (features and labels)
- ▶ In RL:
 - ▶ There is no supervisor: the agent learns from its own experience as it explores
 - ▶ There is no dataset!
- ▶ This leads to the exploration vs exploitation dilemma

EXPLORATION VS. EXPLOITATION DILEMMA

- ▶ Making a decision involves a fundamental choice:
 - ▶ Exploitation: Make the best decision given current information
 - ▶ Exploration: Gather more information
- ▶ The best long-term strategy may involve short-term sacrifices
- ▶ Gather enough information to make the best overall decisions
- ▶ E.g. Restaurant selection
 - ▶ Exploitation: Go to your favourite restaurant
 - ▶ Exploration: Try a new restaurant

REINFORCEMENT LEARNING




source: <https://github.com/brianfarris/RLtalk/blob/master/RLtalk.ipynb>

EXAMPLES

- ▶ You send a user an e-mail
 - ▶ User clicks on the link you get $r = 1$
 - ▶ User fails to click on the link after 3 days $r = 0$
- ▶ Playing games
 - ▶ What is the next best action to take in Chess?
 - ▶ Chess has a sequential element - hence “Reinforcement Learning”
 - ▶ But close enough...
- ▶ Online adverts
 - ▶ User clicks on an advert ($r = 1$)
 - ▶ User fails to click on an advert ($r = 0$)

MULTI-ARMED BANDITS



source: Microsoft Research

- ▶ Agent
- ▶ Action
- ▶ Reward
- ▶ One single state

THE MULTI-ARMED BANDIT PROBLEM



- ▶ A bandit is a tuple $\langle A, R \rangle$
- ▶ Where $a \in A$ is an action out of a set of actions (or “arms”)
- ▶ $r \in R$ is a reward from a set of rewards
- ▶ $R(a, r) = P(r|a)$
 - ▶ The probability of getting a reward r given that I have done action a
 - ▶ It's an unknown probability distribution over rewards
- ▶ At each step, the agent selects an action
- ▶ “You do an action, you get a reward”

THE GOAL

- ▶ Find an optimal policy $\pi(a) = P(a)$ that maximises the long term sum of rewards

- ▶ Long term cumulative reward is $\sum_{t=0}^T r_t$

- ▶ The *action-value* function is the expected reward for taking action a

- ▶ $Q(a) = E[r|a]$

- ▶ The *value* function is $V = E_{\pi}[r]$

- ▶ The expected reward, given the policy I'm following
 - ▶ Optimal $V^* = Q(a^*) = \max_{a \in A} Q(a)$

EXAMPLE

- Three actions to choose from

```
def action_0():  
    return np.random.choice([1, 0], p=[0.5, 0.5])  
  
def action_1():  
    return np.random.choice([1, 0], p=[0.6, 0.4])  
  
def action_2():  
    return np.random.choice([1, 0], p=[0.2, 0.8])  
  
rewards = [action_0, action_1, action_2]  
  
print(rewards[0]()) # 0  
print(rewards[0]()) # 1  
print(rewards[0]()) # 0  
print(rewards[0]()) # 0
```

LET'S SIMULATE: $Q(a_i)$

```
In [32]: ▶ def action_0():  
           return np.random.choice([1,0], p=[0.5, 0.5])  
  
           def action_1():  
               return np.random.choice([1,0], p=[0.6, 0.4])  
  
           def action_2():  
               return np.random.choice([1,0], p=[0.2, 0.8])  
  
           rewards = [action_0, action_1, action_2]
```

```
In [35]: ▶ pulls = 100000  
  
           action_value = []  
           #for action in range(len(rewards)):  
           #    value = [rewards[action]() for _ in range(pulls)]  
           #    action_value.append(value)  
           for reward in rewards:  
               value = [reward() for _ in range(pulls)]  
               action_value.append(value)
```

```
In [36]: ▶ for action, value in enumerate(action_value):  
           print("Action %d: Q(a_%d)=%.2f" % (action, action, np.mean(value)))
```

```
Action 0: Q(a_0)=0.50  
Action 1: Q(a_1)=0.60  
Action 2: Q(a_2)=0.20
```

LET'S SIMULATE: V (1)

```
In [50]: M p0, p1, p2 = 0.33, 0.33, 0.34
```

```
def policy():  
    return np.random.choice([0, 1, 2], p=[p0, p1, p2])
```

```
In [51]: M tot_reward = 0  
for pull in range(pulls):  
    action = policy()  
    tot_reward += rewards[action]()  
print("Total reward =", tot_reward)  
print("Average reward: V =", tot_reward/pulls)
```

```
Total reward = 43000  
Average reward: V = 0.43
```

```
In [52]: M # Manually:  
V = np.mean(action_value[0])*p0 + np.mean(action_value[1]) * p1 + np.mean(action_value[2]) * p2  
print("V =", V)
```

```
V = 0.4308734
```

```
In [53]: M # With the formula:  
V = 0.5 * p0 + 0.6 * p1 + 0.2 * p2  
print("V =", V)
```

```
V = 0.431
```

LET'S SIMULATE: V (2)

```
In [54]: M p0, p1, p2 = 0.4, 0.5, 0.1

def policy():
    return np.random.choice([0, 1, 2], p=[p0, p1, p2])
```

```
In [55]: M tot_reward = 0
for pull in range(pulls):
    action = policy()
    tot_reward += rewards[action]()
print("Total reward =", tot_reward)
print("Average reward: V =", tot_reward/pulls)

Total reward = 51840
Average reward: V = 0.5184
```

```
In [56]: M # Manually:
V = np.mean(action_value[0])*p0 + np.mean(action_value[1]) * p1 + np.mean(action_value[2]) * p2
print("V =", V)

V = 0.519739
```


```
In [57]: M # With the formula:
V = 0.5 * p0 + 0.6 * p1 + 0.2 * p2
print("V =", V)

V = 0.52
```

GOALS (1)

- ▶ So our goal is to find the best action
- ▶ Optimal $V^* = \max_{a \in A} Q(a)$
- ▶ But these values can only be found through averages
 - ▶ $\hat{Q}(a), \hat{V}$
- ▶ We could have done hypothesis testing (recall Lecture 2)
 - ▶ But this would entail a random policy
 - ▶ Maybe we can do better

$$\hat{Q}(a)$$



GOALS (2)

- ▶ We would like to find the best action using the minimum amount of trials possible
- ▶ Keep focusing on the best action
 - ▶ While also checking making sure that other actions are sufficiently explored
- ▶ This is known as the “exploration/exploitation” dilemma

REGRET

- ▶ *Regret* is the opportunity loss for one step
 - ▶ i.e., the difference between the actual payoff and the one you would have if you had played the best option
 - ▶ $I_t = E[(V^* - Q(a_t))]$
- ▶ *Total regret* is the total opportunity loss
 - ▶ $L_t = E \left[\sum_{t=0}^T (V^* - Q(a_t)) \right] = E \left[\sum_{t=0}^T \left(\max_{a \in A} Q(a) - Q(a_t) \right) \right]$
- ▶ It helps us understand how well an algorithm could possibly do, independently of the scale of the rewards

LET'S SIMULATE: REGRET

```
In [72]: # Regret
V_star = max([np.mean(value) for value in action_value]) # 0.6

tot_regret = 0
for pull in range(pulls):
    tot_regret += (V_star - rewards[policy()]())
print("Regret: I_t = %.2f" % (tot_regret/pulls))

Regret: I_t = 0.08

In [73]: I = (V_star - 0.5) * p0 + (V_star - 0.6) * p1 + (V_star - 0.2) * p2
print("I_t = %.2f" % I)

I_t = 0.08

In [74]: print("Total regret: L_t = %.2f" % tot_regret) # also called cumulative regret

Total regret: L_t = 7896.00
```

COUNTING REGRET

- ▶ The *count* $N_t(a)$ is the number of times we took action a until time t
- ▶ The *gap* Δ_a is the difference between the value of the optimal action and that of the action taken, $\Delta_a = V^* - Q(a)$
- ▶ It turns out that regret can be written in terms of gaps and counts:


$$\blacktriangleright L_t = \sum_{a \in A} (E[N_t(a)\Delta_a])$$

- ▶ A good algorithm ensures small counts for large gaps
- ▶ But we have no clue what the gaps are...

PURE EXPLORATION

- ▶ Somewhat similiar to the A/B case
- ▶ You send more or less the equal number of e-mails
- ▶ Very simple setup
- ▶ Link: [When to Run Bandit Tests Instead of A/B Tests](#)
- ▶ Link: [Split testing vs Multi-Armed Bandits](#)

REGRET OF PURE EXPLORATION




- If an algorithm **always** explores, it will have linear total regret

GREEDY

- ▶ Pure exploitation
- ▶ You always choose the action with the highest $\hat{Q}(a)$
- ▶ Can you see a problem with this?
- ▶ Let's try it out

REGRET OF GREEDY




- ▶ If an algorithm **always** explores, it will have linear total regret
- ▶ If an algorithm **never** explores, it will have linear total regret

ϵ -GREEDY

- ▶ You set a small probability ϵ with which you act randomly
- ▶ The rest of the time $(1 - \epsilon)$ you choose the best action
- ▶ This is a very common (but inefficient) setup
- ▶ What is the optimal ϵ ?

REGRET OF ϵ -GREEDY




- If ϵ is constant, ϵ -greedy has asymptotic linear total regret

DECAYING ϵ -GREEDY

- ▶ Same as ϵ -greedy, but now you decrease ϵ as you choose actions
- ▶ E.g.: We do


```
e *= 0.99
```

REGRET OF DECAYING ϵ -GREEDY




- Decaying ϵ -greedy has logarithmic asymptotic total regret.

OPTIMISM IN THE FACE OF UNCERTAINTY (1)



- ▶ You should try actions with highly uncertain outcomes
 - ▶ You believe the best action is the one you haven't explored enough
 - ▶ Choose the one that has the most potential to be best

OPTIMISM IN THE FACE OF UNCERTAINTY (2)



- ▶ We're more certain about a_2
- ▶ More likely to pick another action
- ▶ Until we find the best one, a^*


UPPER CONFIDENCE BOUNDS (1)

- ▶ So far we've estimated $\hat{Q}(a)$, the expected reward for a given action
- ▶ We will now add something that characterises how big the tail of the distribution is, $U(a)$, so that $Q(a) \leq \hat{Q}(a) + U(a)$
- ▶ Then, we can pick the action with the highest *UCB*
 - ▶ $a^* = \arg \max_{a \in A} \left(\hat{Q}(a) + U(a) \right)$


UPPER CONFIDENCE BOUNDS (2)

- ▶ $UCB1(a) = \hat{Q}(a) + C\sqrt{\frac{\log(t)}{N_t(a)}}$
- ▶ $N_t(a)$ is the times action a was executed
- ▶ t is the current timepoint/time
- ▶ $C \in [0, \infty]$ is a constant — I set it to 0.5 for the plots below
 - ▶ Can you guess what the effect of C is?

REGRET OF UPPER CONFIDENCE BOUNDS



INFLUENCE OF C IN THE REGRET OF UCB ALGORITHM




BOOTSTRAP THOMPSON SAMPLING

- ▶ What if you could take bootstrap samples of the action rewards that we have collected?
- ▶ You would have incorporated the uncertainty within your bootstrap samples
- ▶ If you have a large number of bootstrap samples, you have a distribution over possible $\hat{Q}(a)$
- ▶ Sample from this distribution
- ▶ This is a version of **probability matching** (i.e., selecting an action according to the probability that it is the optimal action)
 - ▶ $\pi(a) = P[\hat{Q}(a) > \hat{Q}(a'), \quad \forall a' \in A]$

PRIORS

- ▶ You can get stuck here as well in a sub-optimal action (like greedy)
- ▶ Add some pseudo-rewards
- ▶ Or act randomly a bit

REGRET OF BOOTSTRAP THOMSON SAMPLING



THE SWITCHING BANDIT PROBLEM

- ▶ What if the rewards change?
- ▶ Because people are bored of your e-mails
 - ▶ They talk to each other
 - ▶ Out of fashion
- ▶ You might want to have continuous adaptation
- ▶ Keeping all values and finding $\hat{Q}(a)$ is expensive
 - ▶ What happens in e-mail 1000? And e-mail 100K?

INCREMENTAL CALCULATION OF THE MEAN

v_t can be the reward or the sum of rewards you got at different steps

$$\hat{Q}_{t+1}(a) = \hat{Q}_t(a) + \overbrace{\frac{v_t - \hat{Q}_t(a)}{t}}^{\text{Error}}$$

$$\hat{Q}_{t+1}(a) = \hat{Q}_t(a) + \frac{1}{t} \overbrace{\left[v_t - \hat{Q}_t(a) \right]}^{\text{Error}}$$

$$\hat{Q}_{t+1}(a) = \hat{Q}_t(a) + \alpha \left[v_t - \hat{Q}_t(a) \right]$$

INCREMENTAL BOOTSTRAP

Oza, Nikunj C., and Russell, Stuart, “Online bagging and boosting.”
2005 IEEE International Conference on Systems, Man and
Cybernetics, Vol. 3, 2005.

THE SEQUENTIAL CASE

- ▶ What if you are to take a series of actions?
- ▶ Surely your current action depends on your future (and your past) actions
- ▶ Hence there is going to be a change in the distribution of rewards
 - ▶ Induced by the experimenter

EXAMPLE E-MAIL CAMPAIGN

- ▶ First e-mail
 - ▶ “Please buy this product”
- ▶ Second e-mail
 - ▶ “Will you buy the add-on?”
- ▶ Third e-mail
 - ▶ “Let us service your product”
- ▶ You want to maximise your rewards
- ▶ Creates a tree of possible actions

TREE

- ▶ Let's draw the tree of the above example
 - ▶ Three different actions for each “state”
- ▶ What do you observe?

INTRODUCING STATES

- ▶ $s \in S$ can be used to differentiate between different “states”, conditioning π , V and Q values on states
- ▶ $\pi(s, a)$, $V(s)$, $Q(s, a)$
- ▶ e.g., in the example above, we have $Q(\text{firstemail}, \text{emailtypeA})$
- ▶ Let's write the rest of the states, the policies, V and Q -Values


EQUILIBRIA

- ▶ We will discuss (very) briefly the notion of equilibria
 - ▶ Imagine you are putting up large advert banners on your website
 - ▶ They hide content
 - ▶ User can click on the top right corner and quit the banner
- ▶ Where should you put the banner?
- ▶ How often should the banner pop-up?

ADVERSARIAL BANDITS

- ▶ Most bandits we discussed until now assume the environment is indifferent
- ▶ i.e. the user will click in the link if she thinks it is interesting for her to click
- ▶ But quite often, people are annoyed by your efforts - so they will try to “adapt” around you
 - ▶ Close the advert super-fast without thinking
 - ▶ or use an ad blocker!
- ▶ Solution - put the advert in random places
 - ▶ Mixed policies
- ▶ Exp3 - but not now

CONTEXTUAL BANDITS



Source: <https://medium.com/emergent-future/simple-reinforcement-learning-with-tensorflow-part-0-q-learning-with-tables-and-neural-networks-d195264329d0>

The contextual bandit extends the MAB model by making the decision conditional on the state of the environment.

The contextual bandit is a tuple $\langle A, S, R \rangle$

THE CONTEXTUAL BANDIT PROBLEM

- ▶ Repeat:
 1. Learner presented with **context**
 2. Learner chooses an action
 3. Learner observes reward for the chosen action
- ▶ Goal: Learn a policy that maximises our rewards
- ▶ Issues:
 - ▶ Classic exploitation vs explore dilemma
 - ▶ Plus we also need to learn to use the context effectively
 - ▶ Many actions available
 - ▶ May not see the same context twice — need to generalise
 - ▶ Selection bias: explore while exploiting (i.e., trying to maximise reward), so our data will be skewed

RETHINKING STATES

- ▶ We've seen states as black boxes
 - ▶ They can only be enumerated (i.e. $s_0, s_1 \dots$)
- ▶ What if a state could be decomposed into a set of features? (i.e., *context*)
 - ▶ *Context* is information about the user
 - ▶ *sex, age, married, job...*
- ▶ Learner needs a policy to select the best action for the given context
- ▶ Highly reminiscent of supervised learning
 - ▶ We are given features, we would like to predict an outcome
 - ▶ **What is our outcome?**

COMBINING STATES AND ACTIONS

- ▶ So you now have features that you can encode
- ▶ Various encoding strategies
 - ▶ One regressor per action
 - ▶ A single regressor with encoded actions
- ▶ What could be a problem if you don't have separate regressors for each action?

EXAMPLE

- ▶ Our policy maps our context into an action
- ▶ Note that in the MAB we had *policy()*
- ▶ Before learning, we need to assume a form of policy (e.g., decision tree)

E.g.,

```
def policy(sex, age):  
    ''' Returns an action for the given context '''  
    if sex == 'male':  
        return 1  
    elif age >= 45:  
        return 0  
    else:  
        return 2
```

```
action = policy('female', 30)  
print("Reward", rewards[action]())
```

ϵ -GREEDY AND ϵ -DECREASING

- ▶ Set ϵ to some small value
- ▶ Keep decreasing...
- ▶ Very popular because of its simplicity
- ▶ You need to be smart about your decreasing schedule
 - ▶ Possibly set some lower bound

BOOTSTRAP THOMSON SAMPLING

- ▶ Get a bootstrap sample of all your data
- ▶ Learn a regressor
- ▶ Act greedily using the regressor you learned
- ▶ Repeat

CONCLUSION

- ▶ First hit on bandits
- ▶ Super-exciting research area
- ▶ Used quite a bit on website optimisation and recommender systems
- ▶ We will delve deeper in the adversarial case in the future
- ▶ Again, the bootstrap saves the day